

Class 7: Machine Learning 1

Ellice Wang (PID: A16882742)

2025-01-28


Today we will explore unsupervised machine learning methods including clustering & dimensionality reduction methods.

Let's start by making up some data (where we know there are clear groups) that we can use to test out different clustering methods.

We can use 'rnorm()' function to help us here:

```
hist(rnorm(n=3000, mean = 3))
```

Histogram of rnorm(n = 3000, mean = 3)




Make data with two “clusters”

```
x <- c(rnorm(30, mean=-3), rnorm(30, mean=+3))  
  
z <- cbind(x=x, y=rev(x))  
head(z)
```

```
      x      y  
[1,] -3.190996 1.623297  
[2,] -1.989774 2.755709  
[3,] -2.155930 2.440371  
[4,] -1.993812 3.998618  
[5,] -2.200725 3.642545  
[6,] -1.804729 1.871042
```

```
plot(z)
```



K-means clustering

The main function in “base” R for K-means clustering is called ‘kmeans()’

```
k <- kmeans(z, center = 2)
k
```

K-means clustering with 2 clusters of sizes 30, 30

Cluster means:

```
      x      y
1  2.794307 -3.086687
2 -3.086687  2.794307
```

Clustering vector:

```
[1] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 1 1 1 1 1 1 1
[39] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
```

Within cluster sum of squares by cluster:

```
[1] 59.59596 59.59596
(between_SS / total_SS =  89.7 %)
```

Available components:

```
[1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss"
[6] "betweenss"    "size"         "iter"         "ifault"       "
```

```
attributes(k)
```

\$names

```
[1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss"
[6] "betweenss"    "size"         "iter"         "ifault"       "
```

\$class

```
[1] "kmeans"
```

Q. How many points lie in each cluster?

```
k$size
```

```
[1] 30 30
```

Q. What component of our results tells us about the cluster membership (i.e. which point lives in which cluster)?

```
k$cluster
```

```
[1] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 1 1 1 1 1 1 1 1  
[39] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
```


Q. Center of each cluster?

```
k$centers
```

```
      x      y  
1  2.794307 -3.086687  
2 -3.086687  2.794307
```


Q. Put this result info together and make a little “base R” plot of our clustering result. Also add the cluster center points to this plot.

```
plot(z, col=(k$cluster))  
points(k$centers, col="blue", pch=15, cex=2)
```



Q. Run kmeans on our input ‘z’ and define 4 clusters

```
four_cluster <- kmeans(z, centers=4)
plot(z, col=four_cluster$cluster)
points(four_cluster$centers, col="darkgrey", pch=15, cex=1.5)
```



```
# total sum of squares which tells us how spread out it is
# the smaller the totss, the better the visualization
four_cluster$totss
```

```
[1] 1156.775
```

Hierarchical Clustering


The main function in base R for this called ‘hclust()’. It will take as input a distance matrix (key point is that you can’t just give your “raw” data as input - must first calculate a distance matrix from data)

```
d <- dist(z)
hc <- hclust(d)
hc
```

```
Call:  
hclust(d = d)
```

```
Cluster method   : complete  
Distance         : euclidean  
Number of objects: 60
```


```
plot(hc)  
abline(h=10, col="red")
```



Once I inspect the “tree” I can “cut” the tree to yield my groupings or clusters. The function to this is called ‘cutree()’

```
grps<- cutree(hc, h=10)
```

```
plot(z, col=grps)
```



Hands on with Principle Component Analysis (PCA)

Let's examine some 17-dimensional data.

```
# load in dataframe of UK foods
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url)

# explore the data
dim(x)
```

```
[1] 17  5
```

Q1. How many rows and columns are in your new data frame named x? What R functions could you use to answer this questions?

There are 17 rows and 5 columns in the data frame. I used 'dim()' to answer this question, but 'nrow()' and 'ncol()' would also work.

```
# preview first 6 rows
head(x)
```

	X	England	Wales	Scotland	N.Ireland
1	Cheese	105	103	103	66
2	Carcass_meat	245	227	242	267
3	Other_meat	685	803	750	586
4	Fish	147	160	122	93
5	Fats_and_oils	193	235	184	209
6	Sugars	156	175	147	139

```
# removes first column names
```

```
rownames(x) <- x[,1]
x <- x[,-1]
head(x)
```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139

```
# check dimensions of edited dataframe
dim(x)
```

```
[1] 17  4
```

```
# alternative way to edit dataframe
```


```
x <- read.csv(url, row.names=1)
head(x)
```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139

Q2. Which approach to solving the ‘row-names problem’ mentioned above do you prefer and why? Is one approach more robust than another under certain circumstances?

I prefer using the method that sets the ‘row.names’ argument when loading in the dataframe. It is a more streamlined way to input data. I would use it after checking what the data looked like and determining whether it should be used. Setting ‘x <- x[,-1]’ multiple times will remove the first column multiple times.

```
barplot(as.matrix(x), beside=F, col=rainbow(nrow(x)))
```



```
test <- as.matrix(x)
color <- rainbow(nrow(x))
test <- cbind(test, color)
test
```

	England	Wales	Scotland	N.Ireland	color
Cheese	"105"	"103"	"103"	"66"	"#FF0000"
Carcass_meat	"245"	"227"	"242"	"267"	"#FF5A00"
Other_meat	"685"	"803"	"750"	"586"	"#FFB400"
Fish	"147"	"160"	"122"	"93"	"#F0FF00"
Fats_and_oils	"193"	"235"	"184"	"209"	"#96FF00"
Sugars	"156"	"175"	"147"	"139"	"#3CFF00"

Fresh_potatoes	"720"	"874"	"566"	"1033"	"#00FF1E"
Fresh_Veg	"253"	"265"	"171"	"143"	"#00FF78"
Other_Veg	"488"	"570"	"418"	"355"	"#00FFD2"
Processed_potatoes	"198"	"203"	"220"	"187"	"#00D2FF"
Processed_Veg	"360"	"365"	"337"	"334"	"#0078FF"
Fresh_fruit	"1102"	"1137"	"957"	"674"	"#001EFF"
Cereals	"1472"	"1582"	"1462"	"1494"	"#3C00FF"
Beverages	"57"	"73"	"53"	"47"	"#9600FF"
Soft_drinks	"1374"	"1256"	"1572"	"1506"	"#F000FF"
Alcoholic_drinks	"375"	"475"	"458"	"135"	"#FF00B4"
Confectionery	"54"	"64"	"62"	"41"	"#FF005A"


Q3: Changing what optional argument in the above `barplot()` function results in the following plot?

instead of `'beside=T'` I set `'beside=F'` to generate the stacked barplot

Q5: Generating all pairwise plots may help somewhat. Can you make sense of the following code and resulting figure? What does it mean if a given point lies on the diagonal for a given plot?

The result is a matrix of scatterplots of the dataframe comparing the amount of different food groups each country eats. If a given point is on the diagonal for a given plot, it means that the two countries each similar amounts of that food. For example, England vs Wales plots are in the 2x1 and 1x2 positions. The blue dot lies on the diagonal for these two countries, indicating that England and Wales consume similar amounts of the food represented by the blue dot.

```
pairs(x, col=rainbow(10), pch=16)
```



Q6. What is the main differences between N. Ireland and the other countries of the UK in terms of this data-set?

Northern Ireland's biggest difference is their consumption of the foods represented by the orange and blue dots. Orange represents carcass meat and the blue represents fresh fruit.

Looking at these types of "pairwise plots" can be helpful but it does not scale well & kind of sucks...

PCA to the rescue!

The main function for PCA in base R is called 'prcomp()'. This function wants the transpose of our input data - i.e. the important foods in as columns and the countries as rows.

```
# transpose data
head(t(x))
```

	Cheese	Carcass_meat	Other_meat	Fish	Fats_and_oils	Sugars
England	105	245	685	147	193	156
Wales	103	227	803	160	235	175
Scotland	103	242	750	122	184	147
N.Ireland	66	267	586	93	209	139

	Fresh_potatoes	Fresh_Veg	Other_Veg	Processed_potatoes	
England	720	253	488		198
Wales	874	265	570		203
Scotland	566	171	418		220
N.Ireland	1033	143	355		187
	Processed_Veg	Fresh_fruit	Cereals	Beverages	Soft_drinks
England	360	1102	1472	57	1374
Wales	365	1137	1582	73	1256
Scotland	337	957	1462	53	1572
N.Ireland	334	674	1494	47	1506
	Alcoholic_drinks	Confectionery			
England	375	54			
Wales	475	64			
Scotland	458	62			
N.Ireland	135	41			

```
# Use the prcomp() PCA function
pca <- prcomp( t(x) )
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4
Standard deviation	324.1502	212.7478	73.87622	3.176e-14
Proportion of Variance	0.6744	0.2905	0.03503	0.000e+00
Cumulative Proportion	0.6744	0.9650	1.00000	1.000e+00

Let's see what is in our PCA results

```
attributes(pca)
```

```
$names
[1] "sdev"      "rotation" "center"   "scale"    "x"


$class
[1] "prcomp"
```

The 'pca\$x' result object is where we will focus first as this details how the countries are related to each other in terms of our new "axis"(a.k.a. PCs).

```
pca$x[,2]
```

```
England      Wales      Scotland  N.Ireland
-2.532999 -224.646925  286.081786  -58.901862
```

```
plot(pca$x[,1], pca$x[,2], xlab="PC1", ylab="PC2", xlim=c(-270,500))
text(pca$x[,1], pca$x[,2], colnames(x), col=c("orange", "red", "blue",
"darkgreen"))
```




We can look at the so-called PC “loadings” result to see how the original foods contribute to our new PCs (i.e. how the original variables contribute to our new better variables)

```
pca$rotation
```

	PC1	PC2	PC3	PC4
Cheese	-0.056955380	0.016012850	0.02394295	-0.694538519
Carcass_meat	0.047927628	0.013915823	0.06367111	0.489884628
Other_meat	-0.258916658	-0.015331138	-0.55384854	0.279023718
Fish	-0.084414983	-0.050754947	0.03906481	-0.008483145
Fats_and_oils	-0.005193623	-0.095388656	-0.12522257	0.076097502
Sugars	-0.037620983	-0.043021699	-0.03605745	0.034101334

Fresh_potatoes	0.401402060	-0.715017078	-0.20668248	-0.090972715
Fresh_Veg	-0.151849942	-0.144900268	0.21382237	-0.039901917
Other_Veg	-0.243593729	-0.225450923	-0.05332841	0.016719075
Processed_potatoes	-0.026886233	0.042850761	-0.07364902	0.030125166
Processed_Veg	-0.036488269	-0.045451802	0.05289191	-0.013969507
Fresh_fruit	-0.632640898	-0.177740743	0.40012865	0.184072217
Cereals	-0.047702858	-0.212599678	-0.35884921	0.191926714
Beverages	-0.026187756	-0.030560542	-0.04135860	0.004831876
Soft_drinks	0.232244140	0.555124311	-0.16942648	0.103508492
Alcoholic_drinks	-0.463968168	0.113536523	-0.49858320	-0.316290619
Confectionery	-0.029650201	0.005949921	-0.05232164	0.001847469


```
## Lets focus on PC1 as it accounts for > 90% of variance
par(mar=c(10, 3, 0.35, 0))
barplot( pca$rotation[,1], las=2 )
```



Q9: Generate a similar 'loadings plot' for PC2. What two food groups feature prominently and what does PC2 mainly tell us about?

Fresh potatoes and soft drinks feature prominently (similar to PC1). The other food groups feature less prominently than in the first loading graph. PC2 is the vector that is perpendicular to PC1 and tells us about the second largest source of variation in the data not captured by PC1.


```
# loading plot for PC2
par(mar=c(10, 3, 0.35, 0))
barplot( pca$rotation[,2], las=2 )
```



```
library(ggplot2)


df <- as.data.frame(pca$x)
df_lab <- tibble::rownames_to_column(df, "Country")

# Our first basic plot
ggplot(df_lab) +
  aes(PC1, PC2, col=Country) +
  geom_point()
```



Nicer looking ggplot


```
ggplot(df_lab) +  
  aes(PC1, PC2, col=Country, label=Country) +  
  geom_hline(yintercept = 0, col="gray") +  
  geom_vline(xintercept = 0, col="gray") +  
  geom_point(show.legend = FALSE) +  
  geom_label(hjust=1, nudge_x = -10, show.legend = FALSE) +  
  expand_limits(x = c(-300,500)) +  
  xlab("PC1 (67.4%)") +  
  ylab("PC2 (28%)") +  
  theme_bw()
```

loading plots


```
ld <- as.data.frame(pca$rotation)
ld_lab <- tibble::rownames_to_column(ld, "Food")

ggplot(ld_lab) +
  aes(PC1, Food) +
  geom_col()
```




Nicer looking loading plot

```
ggplot(ld_lab) +
  aes(PC1, reorder(Food, PC1), bg=PC1) +
  geom_col() +
  xlab("PC1 Loadings/Contributions") +
  ylab("Food Group") +
  scale_fill_gradient2(low="purple", mid="gray", high="darkgreen", guide=NULL) +
  theme_bw()
```



Biplots are another way to visualize the information.

```
biplot(pca)
```



PCA of RNA-seq data

```
url2 <- "https://tinyurl.com/expression-CSV"
rna.data <- read.csv(url2, row.names=1)
head(rna.data)
```

	wt1	wt2	wt3	wt4	wt5	ko1	ko2	ko3	ko4	ko5
gene1	439	458	408	429	420	90	88	86	90	93
gene2	219	200	204	210	187	427	423	434	433	426
gene3	1006	989	1030	1017	973	252	237	238	226	210
gene4	783	792	829	856	760	849	856	835	885	894
gene5	181	249	204	244	225	277	305	272	270	279
gene6	460	502	491	491	493	612	594	577	618	638

Q10: How many genes and samples are in this data set?
There are 100 genes and 10 samples in this dataset.


```
dim(rna.data)
```

```
[1] 100  10
```

Plot a PCA graph of the RNA-seq data

```
## Again we have to take the transpose of our data
pca <- prcomp(t(rna.data), scale=TRUE)

## Simple un polished plot of pc1 and pc2
plot(pca$x[,1], pca$x[,2], xlab="PC1", ylab="PC2")
```



```
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	9.6237	1.5198	1.05787	1.05203	0.88062	0.82545	0.80111
Proportion of Variance	0.9262	0.0231	0.01119	0.01107	0.00775	0.00681	0.00642
Cumulative Proportion	0.9262	0.9493	0.96045	0.97152	0.97928	0.98609	0.99251

	PC8	PC9	PC10
Standard deviation	0.62065	0.60342	3.457e-15
Proportion of Variance	0.00385	0.00364	0.000e+00
Cumulative Proportion	0.99636	1.00000	1.000e+00

Using ggplot

```
library(ggplot2)

## Variance captured per PC
pca.var <- pca$sdev^2

## Percent variance is often more informative to look at
pca.var.per <- round(pca.var/sum(pca.var)*100, 1)
```

```
df <- as.data.frame(pca$x)

# Add a 'wt' and 'ko' "condition" column to our plot
df$samples <- colnames(rna.data)
df$condition <- substr(colnames(rna.data),1,2)

p <- ggplot(df) +
  aes(PC1, PC2, label=samples, col=condition) +
  geom_label(show.legend = FALSE) +
# add titles and labels and change theme to make graph look nicer
  labs(title="PCA of RNASeq Data",
        subtitle = "PC1 clealy separates wild-type from knock-out samples",
        x=paste0("PC1 (", pca.var.per[1], "%)"),
        y=paste0("PC2 (", pca.var.per[2], "%)"),
        caption="Class example data") +
  theme_bw()

p
```

